

Explicit SABR Calibration through Simple Expansions

Fabien Le Floc'h^{*} and Gary Kennedy[†]

^{*}*Calypso Technology, 106 rue de La Boétie, 75008 Paris*

[†]*Clarus Financial Technology, London*

(v1.0 released June 2014)

ABSTRACT *The SABR stochastic volatility model is a very popular interpolator of implied volatilities, with a given dynamic. This paper presents a simple and very fast method to calibrate the SABR model to given market volatilities, that is to imply the SABR parameters from a given market smile.*

KEY WORDS: stochastic volatility, SABR, calibration, implied volatility, finance

1. Introduction

The SABR stochastic volatility model (Hagan *et al.*, 2002) enjoys a high popularity as interpolator of implied volatilities, because, in practice, it fits rather well market implied volatility smiles for interest rate derivatives, fx derivatives and equity derivatives with a few parameters α, ρ, ν , and the dynamic of it can be easily controlled through its β parameter, usually defined from historical series analysis for the relevant market.

There are however known shortcomings: it is not arbitrage free as the probability density can become negative for low strikes and long maturities. Many authors have proposed various improvements to the original formula (Oblój, 2008; Johnson and Nonas, 2009; Paulot, 2009; Benaim *et al.*, 2008). More recently the focus has been on finite difference techniques to guarantee the arbitrage-free property (Hagan *et al.*, 2014; Le Floc'h and Kennedy, 2014) in a shifted SABR framework, to allow for negative rates.

In practice, the adjustments to the original formula don't matter much in order to find a good initial guess for the calibration procedure. Once a good initial guess is found, a fast local minimizer like Levenberg-Marquardt can be used to calibrate the specific SABR implementation.

A first calibration method was described in (West, 2005), fitting α exactly to the at-the-money implied volatility, and reducing the problem to a two-dimensional minimization in (ρ, ν) . The Nelder-Mead method is proposed as minimizer. It however still requires some initial guess, and can sometimes be unstable, especially with constraints (Le Floc'h, 2012). A numerical method to find an initial guess that fits exactly the at-the-money volatility and the at-the-money skew is proposed in (Gauthier and Rivaille, 2009). It requires however to solve numerically a two-dimensional non linear system of equations.

In contrast, the method proposed here consists in solving analytically a simple system of equations to fit the at-the-money volatility, skew, and curvature. A similar approach was

Correspondence Address: Calypso Technology, 106 rue de La Boétie, 75008 Paris. Email: fabien.lefloch@calypso.com

applied to the Heston model in (Forde *et al.*, 2012) with the difference that, in their case, the Heston parameters describe a full volatility surface and the initial guess procedure relies on two expiries plus the short time at-the-money volatility to solve the 5 Heston parameters exactly. In the SABR case, there is just one expiry to consider and 3 parameters to fit. We will first present our method for the lognormal formula as well as for the normal formula, and then show its accuracy in calibrating real world smiles on the arbitrage-free model from Hagan *et al.* (2014) compared to a global optimization via differential evolution.

2. SABR normal and lognormal formulas

2.1 The SABR model

In the shifted SABR model, the asset forward follows the following stochastic equation:

$$dF = V(F + b)^\beta dW_1 \quad (1)$$

$$dV = \nu V dW_2 \quad (2)$$

with W_1, W_2 Brownian motions correlated with correlation ρ , and $F(0) = f$, $V(0) = \alpha$. ν represents the volatility of volatility, b is a displacement allowing for negative rates ($b = 0$ for the classic SABR model).

2.2 Normal formula

Given the low rates environment we currently experience, it is now common practice to quote swaptions in terms of normal volatility (bpvol). Taking into account the remarks from Oblój (2008) for the case of the lognormal formula to ensure consistency when $\beta \rightarrow 1$, the expansion of the normal volatility adapted from Hagan *et al.* (2002) for the shifted SABR model is:

for $f \neq K$ and $\beta \in [0, 1]$

$$\sigma_N(K) = \frac{f - K}{x(K)} \left[1 + \left(g(K) + \frac{1}{4} \rho \nu \alpha \beta (f + b)^{\frac{\beta-1}{2}} (K + b)^{\frac{\beta-1}{2}} + \frac{1}{24} (2 - 3\rho^2) \nu^2 \right) T \right] \quad (3)$$

with

$$\begin{aligned} g(K) &= \frac{1}{24} (\beta^2 - 2\beta) (f + b)^{\beta-1} (K + b)^{\beta-1} \alpha^2 \\ \zeta(K) &= \frac{\nu}{\alpha(1-\beta)} \left((f + b)^{1-\beta} - (K + b)^{1-\beta} \right) \\ x(K) &= \frac{1}{\nu} \log \left(\frac{\sqrt{1 - 2\rho\zeta(K) + \zeta^2(K)} - \rho + \zeta(K)}{1 - \rho} \right) \end{aligned}$$

When $f = K$,

$$\sigma_N(f) = \alpha (f + b)^\beta \left[1 + \left(g(f) + \frac{1}{4} \rho \nu \alpha \beta (f + b)^{\beta-1} + \frac{1}{24} (2 - 3\rho^2) \nu^2 \right) T \right] \quad (4)$$

When $\beta = 1$,

$$g(K) = -\frac{1}{24} \alpha^2, \quad \zeta(K) = \frac{\nu}{\alpha} \log \left(\frac{f + b}{K + b} \right)$$

When $\beta = 0$,

$$g(K) = 0, \quad \zeta(K) = \frac{\nu}{\alpha} (f - K)$$

2.3 Lognormal formula

The lognormal formula is very similar to the normal formula, $f - K$ becomes $\log(\frac{f}{K})$ and g includes a small adjustment:

for $f \neq K$ and $\beta \in [0, 1]$,

$$\sigma_B(K) = \frac{1}{x(K)} \log\left(\frac{f+b}{K+b}\right) \left[1 + \left(g(K) + \frac{1}{4} \rho \nu \alpha \beta (f+b)^{\frac{\beta-1}{2}} (K+b)^{\frac{\beta-1}{2}} + \frac{1}{24} (2-3\rho^2) \nu^2 \right) T \right] \quad (5)$$

with

$$\begin{aligned} g(K) &= \frac{1}{24} (\beta-1)^2 (f+b)^{\beta-1} (K+b)^{\beta-1} \alpha^2 \\ \zeta(K) &= \frac{\nu}{\alpha(1-\beta)} \left((f+b)^{1-\beta} - (K+b)^{1-\beta} \right) \\ x(K) &= \frac{1}{\nu} \log \left(\frac{\sqrt{1-2\rho\zeta(K)+\zeta^2(K)} - \rho + \zeta(K)}{1-\rho} \right) \end{aligned}$$

When $f = K$,

$$\sigma_B(f) = \alpha(f+b)^{\beta-1} \left[1 + \left(g(f) + \frac{1}{4} \rho \nu \alpha \beta (f+b)^{\beta-1} + \frac{1}{24} (2-3\rho^2) \nu^2 \right) T \right] \quad (6)$$

When $\beta = 1$,

$$g(K) = 0 \quad , \quad \zeta(K) = \frac{\nu}{\alpha} \log \left(\frac{f+b}{K+b} \right)$$

When $\beta = 0$,

$$g(K) = \frac{1}{24(f+b)(K+b)} \alpha^2 \quad , \quad \zeta(K) = \frac{\nu}{\alpha} (f-K)$$

It is important not to forget that Hagan derived the lognormal formula as an approximation of the normal formula (Hagan *et al.*, 2002): it won't be accurate or realistic at all in cases where the implied volatility is very high (see Table 1). This is not so realistic for market implied volatilities, but it can have an influence in the global minimization procedure, as global minimizers like differential evolution will try out extreme values.

Of course, in order to obtain a lognormal volatility, it is also possible to solve with high accuracy for the Black implied volatility of a given normal volatility: we can firstly compute corresponding Bachelier option price with a fast evaluation method (see appendix) and secondly use a robust and accurate Black implied volatility solver (Jäckel, 2013; Li and Lee, 2011).

Table 1.: Black volatility obtained from the direct lognormal approximation compared the one solved from the normal approximation using
 $\alpha = 3.24, \beta = 1.0, \rho = -0.998, \nu = 1.69, f = 2014, K = f, T = 0.48$.

Formula	Black volatility	Option price
Lognormal	0.9325	510.19
Normal	0.2526	140.41

For performance reasons, we prefer to rely on the normal formula if the calibration is for bpvol and the lognormal formula if the calibration is for Black volatilities.

2.4 Latest Hagan normal formula

3. Explicit initial guess

3.1 Lognormal volatility

The idea is to find α, ρ, ν that matches the at-the-money implied volatility σ_0 , skew σ'_0 and curvature σ''_0 . A direct application of Hagan formula would lead to a system of trivariate polynomials of degrees 3 and 4, which would then require a numerical method. Instead, we rely a simple lower order expansion of the lognormal formula in the variable $z = \log\left(\frac{K+b}{f+b}\right)$ around $z = 0$ that still captures the smile at-the-money accurately (see Appendix XXX):

$$\begin{aligned} \sigma_B(z) &= \alpha(f+b)^{\beta-1} + \frac{1}{2} \left(\rho\nu - (1-\beta)\alpha(f+b)^{\beta-1} \right) z \\ &\quad + \frac{1}{12\alpha(f+b)^{\beta-1}} \left((1-\beta)^2(\alpha(f+b)^{\beta-1})^2 + \nu^2(2-3\rho^2) \right) z^2 \end{aligned} \quad (7)$$

This is the same expansion used in (Hagan *et al.*, 2002, equation (3.1a)) to explain the phenomenology and dynamic of the SABR model. We have then:

$$\begin{cases} \sigma_0 &= \alpha(f+b)^{\beta-1} \\ \sigma'_0 &= \frac{1}{2}(\rho\nu - (1-\beta)\sigma_0) \\ \sigma''_0 &= \frac{1}{3\sigma_0}\nu^2 + \frac{1}{6\sigma_0}((1-\beta)^2\sigma_0^2 - 3\rho^2\nu^2) \end{cases} \quad (8)$$

Note that the derivatives are expressed in log-moneyness towards the variable $z = \log\left(\frac{K+b}{f+b}\right)$ where K is the strike and f is the forward. This system can be exactly solved to give a first guess α_0, ρ_0, ν_0 :

$$\begin{cases} \alpha_0 &= \sigma_0(f+b)^{1-\beta} \\ \nu_0^2 &= 3\sigma_0\sigma''_0 - \frac{1}{2}(1-\beta)^2\sigma_0^2 + \frac{3}{2}(2\sigma'_0 + (1-\beta)\sigma_0)^2 \\ \rho_0 &= \frac{1}{\nu_0}(2\sigma'_0 + (1-\beta)\sigma_0) \end{cases} \quad (9)$$

We make sure to constrain ν_0 to be strictly positive (for example by flooring it to 10^{-4}), and ρ_0 to be in $[-1,1]$. We then refine this first guess by solving exactly for the at-the-money volatility with the given ρ_0 and ν_0 . α_1 is the root of the following third order polynomial:

$$\frac{(1-\beta)^2 T}{24(f+b)^{2-2\beta}} \alpha^3 + \frac{\rho\beta\nu T}{4(f+b)^{1-\beta}} \alpha^2 + \left(1 + \frac{2-3\rho^2}{24} \nu^2 T \right) \alpha - \sigma_0(f+b)^{1-\beta} = 0 \quad (10)$$

This is the same polynomial used by West (2005) to reduce the dimension of the calibration problem. α_1 is picked as the smallest positive root. In theory, choosing the α_1 closest to the previously calibrated α (either from the previous expiry, or from a previous day) could be better. In practice, it turns out to always be the smallest positive root. Our initial guess is then $(\alpha_1, \rho_0, \nu_0)$.

3.2 Normal volatility

When the input volatilities are normal (bpvol), we can just convert them to equivalent log-normal volatilities (eventually relying on some accurate expansion of the normal volatility at-the-money), and apply the same method as in the previous section to find the initial guess. Instead, we prefer to work directly with normal volatility and derive a simple expansion of the normal formula in $z = \log\left(\frac{K+b}{f+b}\right)$ around $z = 0$ (see Appendix XXX);

$$\sigma_N(z) = \alpha(f+b)^\beta + \frac{1}{2} \left(\rho\nu(f+b) + \beta\alpha(f+b)^\beta \right) z \quad (11)$$

$$+ \left[\frac{1}{12\alpha(f+b)^{\beta-2}} (2\nu^2 - 3\rho^2\nu^2) + \frac{1}{4}\rho\nu(f+b) + \frac{1}{12}(\beta^2 + \beta)\alpha(f+b)^\beta \right] z^2 \quad (12)$$

Let $\sigma_0, \sigma'_0, \sigma''_0$ be respectively the normal volatility, the slope and the curvature at the money expressed in log-moneyness z . We have:

$$\begin{cases} \sigma_0 &= \alpha(f+b)^\beta \\ \sigma'_0 &= \frac{1}{2}\rho\nu(f+b) + \frac{1}{2}\beta\sigma_0 \\ \sigma''_0 &= \frac{(f+b)^2}{6\sigma_0} (2\nu^2 - 3\rho^2\nu^2) + \frac{1}{2}\rho\nu(f+b) + \frac{1}{6}(\beta^2 + \beta)\sigma_0 \end{cases} \quad (13)$$

This system can be exactly solved to give a first guess α_0, ρ_0, ν_0 :

$$\begin{cases} \alpha_0 &= \sigma_0(f+b)^{-\beta} \\ \nu_0^2 &= \frac{1}{(f+b)^2} \left[3\sigma_0\sigma''_0 - \frac{1}{2}(\beta^2 + \beta)\sigma_0^2 - 3\sigma_0(\sigma'_0 - \frac{1}{2}\beta\sigma_0) + \frac{3}{2}(2\sigma'_0 - \beta\sigma_0)^2 \right] \\ \rho_0 &= \frac{1}{\nu_0(f+b)} (2\sigma'_0 - \beta\sigma_0) \end{cases} \quad (14)$$

As for the lognormal volatility guess, we refine this first guess by solving exactly for the at-the-money volatility with the given ρ_0 and ν_0 . α_1 is a root of the following third order polynomial:

$$\frac{(\beta^2 - 2\beta)T}{24(f+b)^{2-2\beta}} \alpha^3 + \frac{\rho\beta\nu T}{4(f+b)^{1-\beta}} \alpha^2 + \left(1 + \frac{2-3\rho^2}{24} \nu^2 T \right) \alpha - \sigma_0(f+b)^{-\beta} = 0 \quad (15)$$

Note that this polynomial is slightly different from the lognormal volatility polynomial. We take the α_1 to be the smallest positive root.

3.3 How to find the at-the-money volatility, slope and curvature?

The simplest is to fit a parabola to the three closest points around the forward with coordinates $(z_{-1}, \hat{\sigma}_{-1}), (z_0, \hat{\sigma}_0), (z_1, \hat{\sigma}_1)$. This is equivalent to a 3 points finite difference on a non uniform grid. We have then:

$$\sigma_0 = z_0 z_1 w_{-1} \hat{\sigma}_{-1} + z_{-1} z_1 w_0 \hat{\sigma}_0 + z_{-1} z_0 w_1 \hat{\sigma}_1 \quad (16)$$

$$\sigma'_0 = -(z_0 + z_1) w_{-1} \hat{\sigma}_{-1} - (z_{-1} + z_1) w_0 \hat{\sigma}_0 - (z_{-1} + z_0) w_1 \hat{\sigma}_1 \quad (17)$$

$$\sigma''_0 = 2w_{-1} \hat{\sigma}_{-1} + 2w_0 \hat{\sigma}_0 + 2w_1 \hat{\sigma}_1 \quad (18)$$

with

$$w_{-1} = \frac{1}{(z_{-1} - z_0)(z_{-1} - z_1)} \quad (19)$$

$$w_0 = \frac{1}{(z_0 - z_{-1})(z_0 - z_1)} \quad (20)$$

$$w_1 = \frac{1}{(z_1 - z_{-1})(z_1 - z_0)} \quad (21)$$

In our experience, using higher order approximations like a 5 points finite difference (equivalent to a quartic on 5 points), or a cubic spline does not lead to any visible improvement in accuracy on our problem.

When the input data is noisy, one could resort to a best fit parabola around the five closest points from the forward, which can be reduced to solving simple linear system. When the volatility at the forward is already present in the input data, we make sure the parabola passes exactly through it.

Another approach is to repeat the initial guess procedure with a parabola on different sets of three points further away from the forward and select the guess corresponding to the best overall fit.

In our numerical tests we will select the guess that gives the minimum mean square error in volatilities between the two guesses stemming from the 3-points parabola and the 5-points parabola.

4. Almost exact inversion of SABR smiles

We use the SABR parameters found by calibration of the lognormal SABR model on the S&P500 options in December 2008 as starting point (Table 2). From those we regenerate a discrete set of implied volatilities for 12 strikes and 11 expiries using the lognormal SABR formula with $\beta = 1$. And finally we apply the explicit initial guess procedure to each expiry.

Table 2.: Reference SABR parameters and explicit guess

Expiry years	Reference parameters			Explicit guess			
	α	ρ	ν	α	ρ	ν	vol RMSE
0.058	0.271	-0.345	1.010	0.271	-0.343	1.008	2.64e-04
0.153	0.256	-0.321	0.933	0.256	-0.319	0.934	1.36e-04
0.230	0.256	-0.346	0.820	0.256	-0.344	0.822	1.13e-04
0.479	0.255	-0.370	0.629	0.255	-0.369	0.631	8.25e-05
0.729	0.257	-0.403	0.528	0.257	-0.402	0.528	2.30e-05
1.227	0.260	-0.429	0.448	0.260	-0.429	0.447	6.15e-05
1.726	0.261	-0.440	0.392	0.261	-0.440	0.390	1.17e-04
2.244	0.262	-0.445	0.355	0.262	-0.445	0.352	1.59e-04
2.742	0.262	-0.445	0.329	0.262	-0.445	0.326	1.85e-04
3.241	0.262	-0.447	0.310	0.262	-0.447	0.306	2.13e-04
4.239	0.263	-0.452	0.284	0.263	-0.452	0.279	2.67e-04

The root mean square error in implied volatilities of the fit is always lower than $3 \cdot 10^{-4}$. ρ, ν are recovered with an accuracy higher than $5 \cdot 10^{-3}$ and α is recovered with an accuracy higher than 10^{-4} , independently of the expiry (see Table 2). The refinement α_1 allows to gain one order of magnitude in accuracy over α_0 , we display the various stages of the initial guess procedure on the slice of expiry 0.479 on Figure 1.


5. Real world smiles

5.1 Calibration on the analytic formula

5.1.1 Equity option smiles

To calibrate SABR to equity implied volatility surfaces from December 2012, we minimize the weighted root mean square error in implied volatilities with the Levenberg-Marquardt method (Levenberg, 1944; Marquardt, 1963), placing more weight to volatilities around ATM

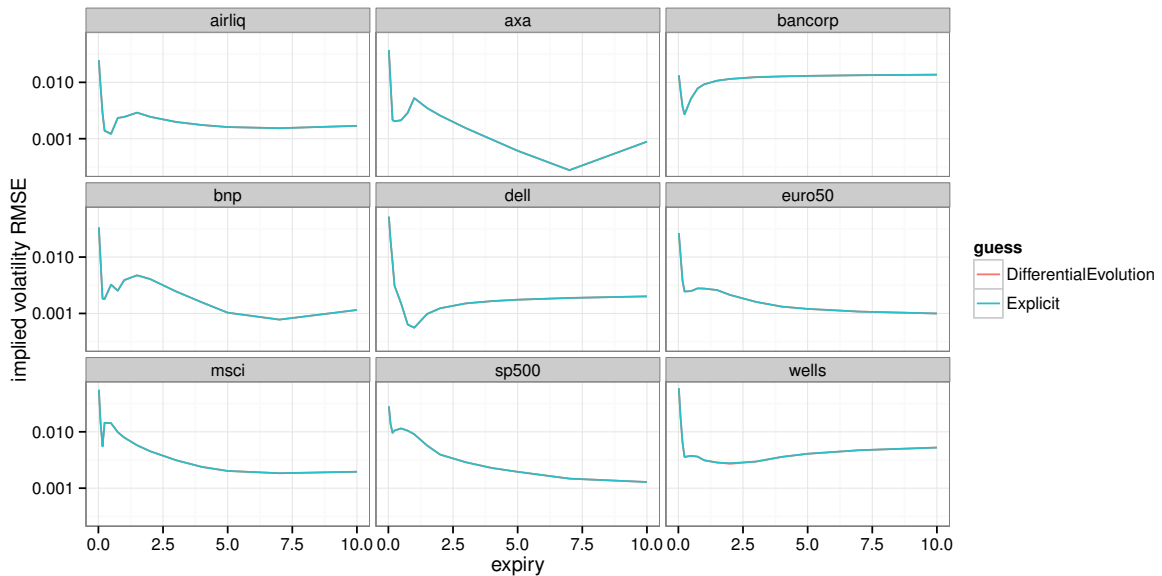

Figure 1.: The various stages of the initial guess method for the slice of expiry 0.479 year



+/- 20%. We compare the results obtained using either our explicit method as initial guess, or alternatively a guess found by running the differential evolution algorithm (Storn and Price, 1997) with a population size 20 on 1000 generations. The resulting mean square error is displayed in Figure 2.

The minimization on the explicit guess leads a mean square error in implied volatilities nearly indistinguishable from a minimization on the differential evolution guess. It is interesting however to look at the calibrated parameters values: the calibrated α shows some strong variations with the guess found by differential evolution, while it is relatively smooth with our explicit initial guess (Figure 3). ρ and ν behave similarly. We will take a closer look at why this happens in the next section.

Figure 2.: Root mean square error of calibration per expiry on various equity surfaces


Figure 3.: Calibrated α on various equity surfaces

5.1.2 Two SABRs for the same smile

Figure 4 illustrates how two different sets of SABR parameters can lead to an equally good fit. We just picked the S&P 500 options expiring in 4 years used in the previous section. The root mean square error is 0.00219269910750194 for the SABR parameters with lower α while the one for the higher α is 0.002192699107502103, that is the lower α fits better by only 10^{-16} .

The fact that two different SABR parameters sets can lead to smiles with nearly the same error measure is the root of issues with differential evolution: it is not obvious which set of parameter is the right one without some additional constraint on the variation of the

Figure 4.: Smiles corresponding to $\alpha = 0.237, \rho = -0.735, \nu = 0.398$ and $\alpha = 0.850, \rho = -0.742, \nu = 1.335$ with $\beta = 1, f = 1426.19, T = 4$



parameters. This can be resolved by some penalty term, but it is difficult in practice to find the correct penalty term that works for a variety of situations. A better approach would be to look beyond the best candidate, and consider as alternative initial guesses the other candidates in the latest generation with very similar error measure and different (lower) α . In contrast, the explicit guess leads directly to the correct smile.

5.1.3 Swaption smiles


We apply the same calibration methodology as for the equity smiles cases, but relying on the normal formula with $\beta = 0.5$ instead. The calibrated parameters are indistinguishable between the calibration on the explicit initial guess and the calibration on the guess found by differential evolution (see Table 3). Furthermore, the initial guess is very close to the calibration result.

Figures 5, 6(a) and 6(b) show how close are the smile produced by the initial guess and the smile resulting from the Levenberg-Marquardt minimization for short expiries as well as for long ones.

Table 3.: calibrated SABR parameters with $\beta = 0.5$. LM denotes Levenberg-Marquardt minimization.

Method	α	ρ	ν	normal vol RMSE	time(ms)
1m5y swaption on May 28 2014					
Explicit	0.052	0.458	0.751	2.95e-04	0.01
Explicit + LM	0.052	0.368	0.768	2.39e-04	0.11
Differential Evolution + LM	0.052	0.368	0.768	2.39e-04	5.17
2y5y swaption on May 28 2014					
Explicit	0.052	0.060	0.322	1.43e-05	0.01
Explicit + LM	0.052	0.058	0.313	7.59e-06	0.10
Differential Evolution + LM	0.052	0.058	0.313	7.59e-06	1.95
10y10y swaption on May 28 2014					
Explicit	0.037	-0.141	0.319	2.26e-05	0.01
Explicit + LM	0.037	-0.153	0.305	4.68e-06	0.12
Differential Evolution + LM	0.037	-0.153	0.305	4.68e-06	3.67

Figure 5.: Initial guess and calibrated smile for a May 2014 1m5y Swaption



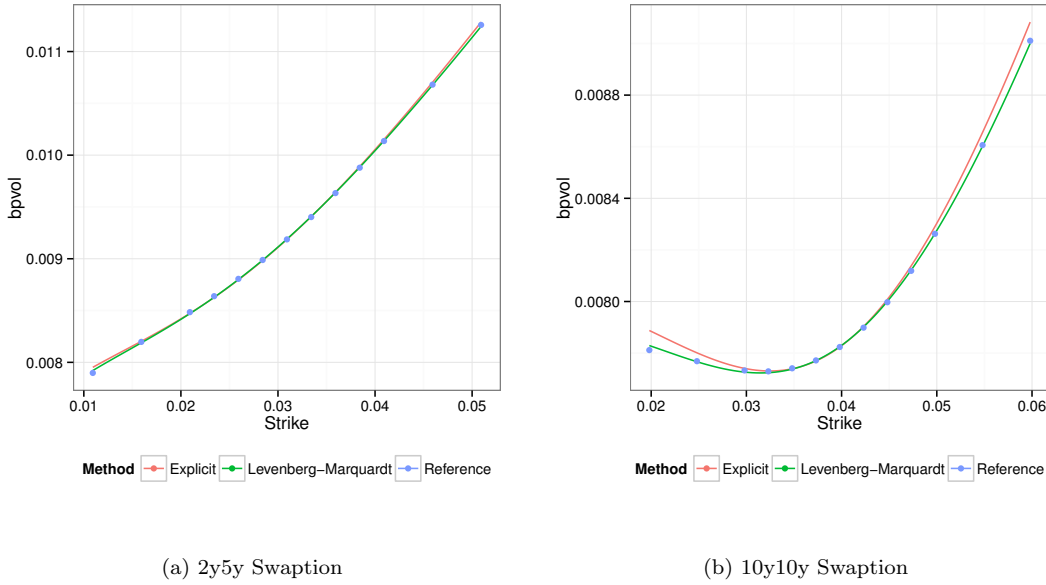


Figure 6.: Initial guess and calibrated smile

5.2 Calibration on the arbitrage free PDE

We calibrate the same swaptions as in the previous section but using the arbitrage free SABR model from (Hagan *et al.*, 2014). In the calibration procedure there is the need to transform the swaption price resulting from integrating the probability density to a normal volatility, in order to compute the error in terms of normal volatilities. This can be achieved with near machine accuracy using the algorithm from (Le Floch, 2014). The transformation to normal volatilities must be reasonably accurate for extremely low option prices, as even if the options reference prices are not so low, the prices resulting from the best fit procedure can be extremely low.

Short and medium expiries calibrated SABR parameters are very close to the one of the previous section (Table 4).

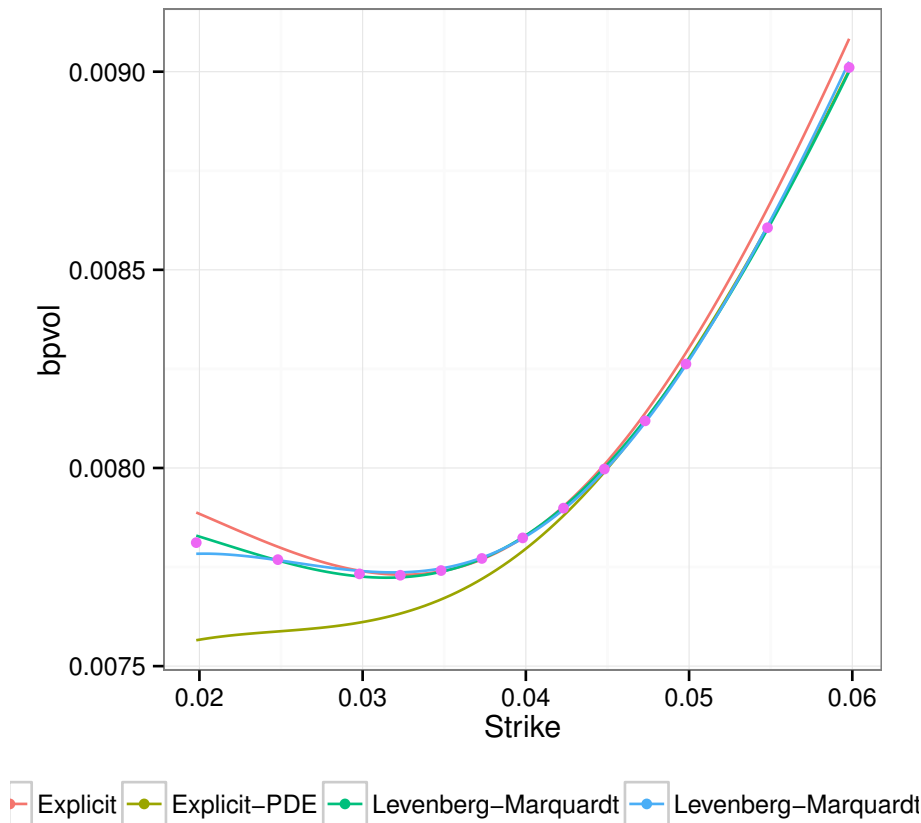
Table 4.: calibrated SABR parameters with $\beta = 0.5$ for the arbitrage free SABR model. LM denotes Levenberg-Marquardt minimization.

Method	α	ρ	ν	normal vol RMSE	time(s)
1m5y swaption on May 28 2014					
Explicit	0.052	0.458	0.751	2.87e-04	0.005
Explicit+LM	0.052	0.374	0.765	2.40e-04	0.071
2y5y swaption on May 28 2014					
Explicit	0.052	0.055	0.328	1.40e-05	0.005
Explicit+LM	0.052	0.058	0.323	6.61e-06	0.081
10y10y swaption on May 28 2014					
Explicit	0.037	-0.145	0.322	7.36e-05	0.005
Explicit+LM	0.037	-0.194	0.359	6.97e-06	0.068

There is however a visible difference for the longest expiry. In particular, the SABR parameters ρ, ν have a slightly different meaning: the smile resulting from the normal explicit initial

guess is lower for low strikes with the arbitrage free PDE approach than the smile computed from the same guess with the normal formula (Figure 7).

Figure 7.: Initial guess and calibrated smile for a May 2014 10y10y Swaption using the Arbitrage-free SABR PDE approach



6. Conclusion

We have presented an explicit initial guess procedure to calibrate the lognormal SABR formula as well as the normal SABR formula, that, despite its simplicity, works particularly well for a variety of different SABR configurations. It allows for faster calibration, by using only a local minimizer, as well as a more reliable calibration, by avoiding the problem of multiple best fits that can appear with global minimizers. We have shown it can also be applied to the arbitrage free SABR model with success, even the guess is further away from the target for longer maturities, due to the mismatch between the arbitrage free SABR model and the classic normal SABR formula.

References

- Benaim, S., Dodgson, M. and Kainth, D., An arbitrage-free method for smile extrapolation. (2008) , Technical report, Working Paper, QuaRC, Royal Bank of Scotland.
- Forde, M., Jacquier, A. and Lee, R. (2012) The small-time smile and term structure of implied volatility under the Heston model, *SIAM Journal on Financial Mathematics*, 3(1), pp. 690–708.
- Gauthier, P. and Rivaille, P. Y. H. (2009) Fitting the smile, Smart parameters for SABR and Heston, *Smart Parameters for SABR and Heston (October 30, 2009)*.
- Hagan, P. S., Kumar, D., Lesniewski, A. S. and Woodward, D. E. (2002) Managing smile risk, *Wilmott magazine*.
- Hagan, P. S., Kumar, D., Lesniewski, A. S. and Woodward, D. E. (2014) Arbitrage free SABR, *Wilmott magazine*.
- Jäckel, P. (2013) Lets be rational. , <http://www.pjaeckel.webspace.virginmedia.com/LetsBeRational.pdf> (accessed 03-April-2014).
- Johnson, S. and Nonas, B. (2009) Arbitrage-free construction of the swaption cube, *Wilmott Journal*, 1(3), pp. 137–143.
- Le Floch, F. (2012) Issues of Nelder-Mead Simplex Optimisation with Constraints, *SSRN eLibrary* <http://papers.ssrn.com/abstract=2097904>.
- Le Floch, F. (2014) Fast and Accurate Analytic Basis Point Volatility, *SSRN eLibrary* <http://papers.ssrn.com/abstract=2420757>.
- Le Floch, F. and Kennedy, G. (2014) Finite difference techniques for arbitrage free SABR, *SSRN eLibrary* <http://papers.ssrn.com/abstract=2402001>.
- Levenberg, K. (1944) A method for the solution of certain problems in least squares, *Quarterly of applied mathematics*, 2, pp. 164–168.
- Li, M. and Lee, K. (2011) An adaptive successive over-relaxation method for computing the Black–Scholes implied volatility, *Quantitative Finance*, 11(8), pp. 1245–1269.
- Marquardt, D. W. (1963) An algorithm for least-squares estimation of nonlinear parameters, *Journal of the Society for Industrial & Applied Mathematics*, 11(2), pp. 431–441.
- Oblój, J. (2008) Fine-tune your smile: Correction to Hagan et al, *Wilmott Magazine*.
- Paulot, L. (2009) Asymptotic implied volatility at the second order with application to the SABR model, *Available at SSRN 1413649*.
- Storn, R. and Price, K. (1997) Differential evolution—a simple and efficient heuristic for global optimization over continuous spaces, *Journal of global optimization*, 11(4), pp. 341–359.
- West, G. (2005) Calibration of the SABR model in illiquid markets, *Applied Mathematical Finance*, 12(4), pp. 371–385.